

Trybe Technical Test - OULAD

Lucas Cardoso de Menezes

07/08/2022

Purpose and context:

The learning team at the University wants to understand what the profile of people students is, what are the factors related to people's performance and what recommendations/initiatives you suggest for people students' performance to improve.

The job was to explore the data available and drive business decision making.

Database: *OULAD - Open University Learning Analytics Dataset*

Description:

A dataset containing demographic information about students, their courses taken, and the final outcomes of each course.

Roadmap

1. More about the dataset;
2. Database schema;
3. Preparing the environment for data analysis;
4. Understanding the profile of the students(I);
5. Performance factors (II);
6. Conclusion(III).

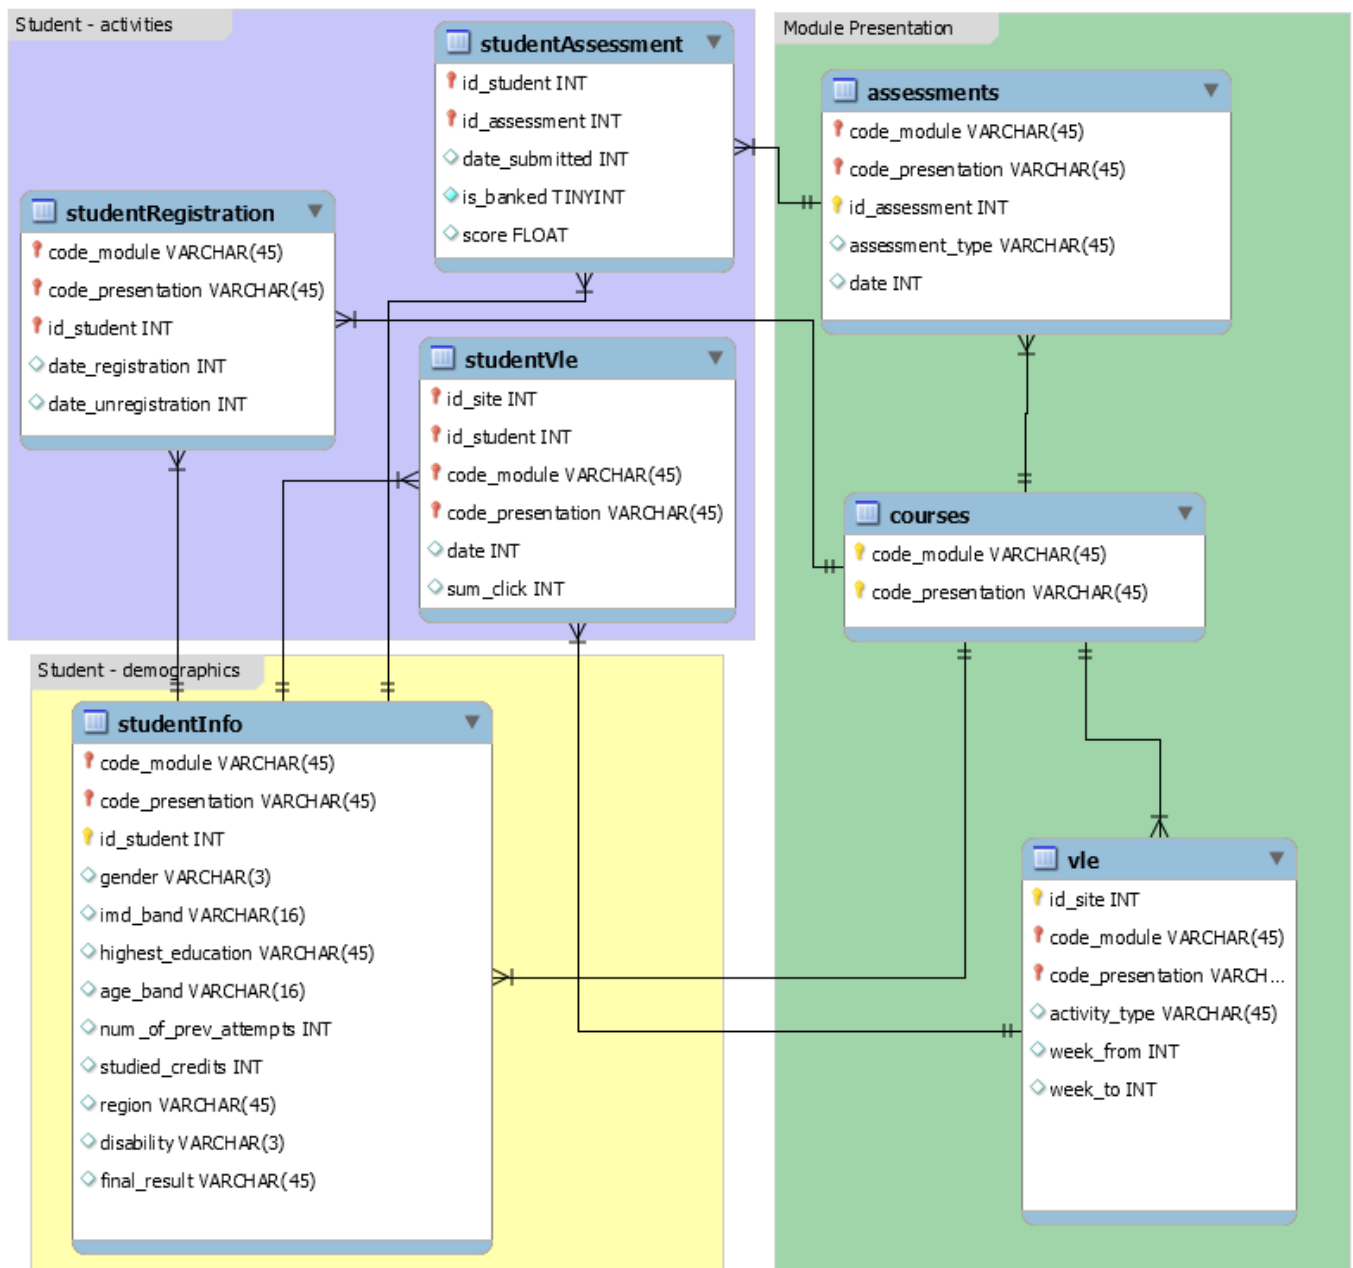
1. More about the dataset

The anonymized *Open University Learning Analytics Dataset (OULAD)*, contains data on courses, students, and their interactions with the *Virtual Learning Environment (VLE)* for seven selected courses. The courses - starting in February and October - are marked as "B" and "J" respectively. The dataset consists of tables connected using unique identifiers. All tables are stored in csv format.

Kuzilek J., Hlosta M., Zdrahal Z. Open University Learning Analytics dataset Sci.

2. database schema

Here we have a schema (https://analyse.kmi.open.ac.uk/open_dataset (https://analyse.kmi.open.ac.uk/open_dataset)) to illustrate the data structure of the dataset.



As you can see, there are many different types of data involved, but since we want to understand the profile of the people students and the performance factors we will use:

- Demographic data of the sample;
- A measure of the students' commitment to the course over the term;
- A measure of their performance over the period.

Going to the indicated website, we can see that this information is contained in the following tables:

- studentInfo;
- studentAssessment;
- assessments;
- studentVle;
- vle.

These tables will be our data sources for meeting the objectives.

Preparing the environment for data analysis

To perform this analysis we will use two R packages, *dplyr* and *plotly*. One to assist in manipulating the tables and the other in generating the graphs that will be presented.

The packages are available at:

- dplyr: <https://cran.r-project.org/web/packages/dplyr/index.html> (<https://cran.r-project.org/web/packages/dplyr/index.html>)
- plotly: <https://cran.r-project.org/web/packages/plotly/index.html> (<https://cran.r-project.org/web/packages/plotly/index.html>)

Or using the commands:

```
# install.packages("dplyr")
# install.packages("plotly")
```

We will use *library* to call the packages after installation:

```
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 4.2.2
```

```
library(plotly)
```

```
## Warning: package 'plotly' was built under R version 4.2.2
```

```
## Warning: package 'ggplot2' was built under R version 4.2.2
```

To finalize the environment preparation we must inform the location where the input data is located:

```
setwd(dir = "C:/Users/luqui/Documents/Trybe")
```

4. understanding the profile of the students(I)

To be able to understand the profile of the students present in the database we will use the **studentInfo** table as it contains demographic information about the students along with their results. The file contains the following columns:

```
student_info <-
  read.csv(
    file = paste0(getwd(), "/Input/studentInfo.csv")
  )

colnames(student_info)
```

```
## [1] "code_module"      "code_presentation" "id_student"
## [4] "gender"           "region"            "highest_education"
## [7] "imd_band"         "age_band"          "num_of_prev_attempts"
## [10] "studied_credits"  "disability"        "final_result"
```

We will use only the columns that contain information linked to the student's demographic profile, and we will also leave only the values without repetition, so that we have unique data for each student:

```
student_info_profile <-  
  student_info[,c(2:6,8,11)] %>%  
  distinct(id_student, .keep_all = TRUE)  
  
colnames(student_info_profile)
```

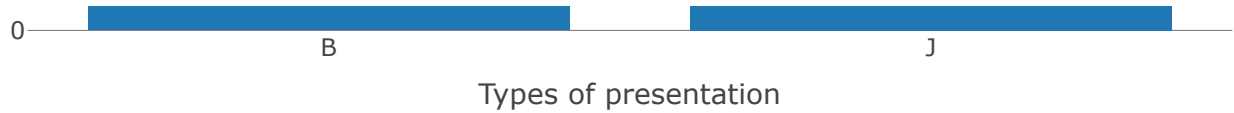
```
## [1] "code_presentation" "id_student"      "gender"  
## [4] "region"            "highest_education" "age_band"  
## [7] "disability"
```

Generating the quantites:

```
#Types of presentation and the number of students  
student_info_presentation <-  
  student_info_profile %>%  
  mutate(  
    type_presentation = substr(code_presentation, nchar(code_presentation), nchar(code_presentation))  
  ) %>%  
  group_by(type_presentation) %>%  
  count() %>%  
  ungroup()  
  
student_info_presentation
```

```
## # A tibble: 2 × 2  
##   type_presentation      n  
##   <chr>              <int>  
## 1 B                11190  
## 2 J                17595
```






#Students per Gender

```
student_info_gender <-  
  student_info_profile %>%  
  group_by(gender) %>%  
  count() %>%  
  ungroup()
```

student_info_gender

A tibble: 2 × 2

```
##   gender      n  
##   <chr>  <int>  
## 1 F      13739  
## 2 M      15046
```



#Students per region

```
student_info_region <-  
  student_info_profile %>%  
  group_by(region) %>%  
  count() %>%  
  ungroup()
```


student_info_region

```
## # A tibble: 13 × 2
##   region      n
##   <chr>    <int>
## 1 East Anglian Region  3000
## 2 East Midlands Region 2095
## 3 Ireland            1072
## 4 London Region       2845
## 5 North Region        1588
## 6 North Western Region 2548
## 7 Scotland            2934
## 8 South East Region    1875
## 9 South Region         2737
## 10 South West Region   2154
## 11 Wales              1876
## 12 West Midlands Region 2269
## 13 Yorkshire Region    1792
```

```
#Education
student_info_education <-
  student_info_profile %>%
  group_by(highest_education) %>%
  count() %>%
  ungroup()

student_info_education
```

```
## # A tibble: 5 × 2
##   highest_education      n
##   <chr>              <int>
## 1 A Level or Equivalent 12355
## 2 HE Qualification      4092
## 3 Lower Than A Level    11780
## 4 No Formal quals        306
## 5 Post Graduate Qualification 252
```



u

A Level or Equivalent

HE Qualification

Lower Than A Level

No Formal quals


Post Graduate Qualification

Education

```
#Age
student_info_age <-
  student_info_profile %>%
  group_by(age_band) %>%
  count() %>%
  ungroup()

student_info_age
```

```
## # A tibble: 3 × 2
##   age_band     n
##   <chr>   <int>
## 1 0-35    20145
## 2 35-55     8462
## 3 55<=      178
```




```
#Disabled
```

```
student_info_disability <-  
  student_info_profile %>%  
  group_by(disability) %>%  
  count() %>%  
  ungroup()
```

```
student_info_disability
```

```
## # A tibble: 2 × 2
```

```
##   disability     n  
##   <chr>       <int>  
## 1 N         26068  
## 2 Y          2717
```



```
#Demographic information compiled
```

```
student_demographic_data <-  
  student_info_profile %>%  
  group_by(gender, region, highest_education, age_band, disability) %>%  
  count() %>%  
  ungroup()
```

```
student_demographic_data
```



```
## # A tibble: 336 × 6
##   gender region          highest_education age_band disability     n
##   <chr>  <chr>          <chr>          <chr>    <chr>    <int>
## 1 F      East Anglian Region A Level or Equivalent 0-35    N      392
## 2 F      East Anglian Region A Level or Equivalent 0-35    Y      70
## 3 F      East Anglian Region A Level or Equivalent 35-55    N     169
## 4 F      East Anglian Region A Level or Equivalent 35-55    Y      29
## 5 F      East Anglian Region HE Qualification      0-35    N      65
## 6 F      East Anglian Region HE Qualification      35-55    N      66
## 7 F      East Anglian Region Lower Than A Level    0-35    N     399
## 8 F      East Anglian Region Lower Than A Level    0-35    Y      60
## 9 F      East Anglian Region Lower Than A Level    35-55    N     200
## 10 F     East Anglian Region Lower Than A Level    35-55    Y      41
## # ... with 326 more rows
```

Observing the graphs, we conclude that we have a multicultural profile of students, coming from different regions, with different levels of knowledge, with the majority being up to 35 years old. However, we decided to correlate the data we consider most important (age, education, and region) to better understand the profile of the sample:

```
#Schooling by student's region
student_info_region_education <-
  student_info_profile %>%
  group_by(region, highest_education) %>%
  count() %>%
  ungroup()

student_info_region_education
```

```
## # A tibble: 60 × 3
##   region          highest_education     n
##   <chr>          <chr>          <int>
## 1 East Anglian Region A Level or Equivalent 1305
## 2 East Anglian Region HE Qualification      313
## 3 East Anglian Region Lower Than A Level    1324
## 4 East Anglian Region No Formal quals       49
## 5 East Anglian Region Post Graduate Qualification 9
## 6 East Midlands Region A Level or Equivalent 944
## 7 East Midlands Region HE Qualification      176
## 8 East Midlands Region Lower Than A Level    960
## 9 East Midlands Region No Formal quals       11
## 10 East Midlands Region Post Graduate Qualification 4
## # ... with 50 more rows
```

When we correlate schooling by region, we see a still homogeneous picture.

```
#Schooling by student age
student_info_age_education <-
  student_info_profile %>%
  group_by(age_band, highest_education) %>%
  count() %>%
  ungroup()

student_info_age_education
```

```
## # A tibble: 14 × 3
##   age_band highest_education      n
##   <chr>    <chr>            <int>
## 1 0-35    A Level or Equivalent    9290
## 2 0-35    HE Qualification        2228
## 3 0-35    Lower Than A Level      8284
## 4 0-35    No Formal quals         258
## 5 0-35    Post Graduate Qualification 85
## 6 35-55   A Level or Equivalent    3032
## 7 35-55   HE Qualification        1756
## 8 35-55   Lower Than A Level      3469
## 9 35-55   No Formal quals         48
## 10 35-55   Post Graduate Qualification 157
## 11 55<=   A Level or Equivalent     33
## 12 55<=   HE Qualification         108
## 13 55<=   Lower Than A Level       27
## 14 55<=   Post Graduate Qualification 10
```

Now, correlating *Age x Education*, we can identify points where there is a larger sample size, so we decided to create a table of only the students that contain this profile:

```
#Schooling by student age
representative_student_group_info <-
  student_info_profile %>%
  subset(
    age_band == "0-35" &
    (
      highest_education == "A Level or Equivalent" |
      highest_education == "Lower Than A Level" |
      highest_education == "HE Qualification"
    )
  )

#How much this group represents
group_percentage <-
  (nrow(representative_student_group_info)/nrow(student_info_profile))*100


group_percentage
```

```
## [1] 68.79277
```

Representative Samples



■ Representative Group
■ Non-Representative Group



After these correlations, it can be seen that those students aged '0-35' who have College Level, or High School complete/studying, represent the general profile of the students, as they are 68.8% of the samples.

5. performance factors (II)

Performance on each assessment is a good indicator of students' knowledge of the course. We will separate the final exams from the other assessments, as their status and participation in the final assessment are different from the others.

Reading the assessment data

```
#Information from the tests per student
student_assessment <-
  read.csv(
    file = paste0(getwd(), "/Input/studentAssessment.csv")
  )

#Tasting information
assessments <-
  read.csv(
    file = paste0(getwd(), "/Input/assessments.csv")
  )
```

Separating the Exams

```
final_exams <-
  assessments %>%
  subset(assessment_type == "Exam")

head(final_exams)
```

```
##      code_module code_presentation id_assessment assessment_type date weight
## 6          AAA      2013J          1757          Exam    NA      100
## 12         AAA      2014J          1763          Exam    NA      100
## 24         BBB      2013B          14990          Exam    NA      100
## 36         BBB      2013J          15002          Exam    NA      100
## 48         BBB      2014B          15014          Exam    NA      100
## 54         BBB      2014J          15025          Exam    NA      100
```

```
others_exams <-
  assessments %>%
  subset(assessment_type != "Exam")

head(others_exams)
```

```
##      code_module code_presentation id_assessment assessment_type date weight
## 1          AAA      2013J          1752          TMA     19      10
## 2          AAA      2013J          1753          TMA     54      20
## 3          AAA      2013J          1754          TMA    117      20
## 4          AAA      2013J          1755          TMA    166      20
## 5          AAA      2013J          1756          TMA    215      30
## 7          AAA      2014J          1758          TMA     19      10
```

Let's identify the average rating per student per module, and identify the activities of those with the highest and lowest average ratings.

```
#Creating the data frame 'student_group_kpis

student_group_kpis <-
  student_assessment %>%
  mutate(pass = ifelse(score>=40, TRUE, FALSE))

#Putting together the exam information and creating the columns of who passed the exam and the grid weight

student_group_others_exams <-
  student_group_kpis %>%
  inner_join(others_exams, by = "id_assessment")

student_group_others_exams <-
  student_group_others_exams %>%
  mutate(weight_grade = score*weight/100)

head(student_group_others_exams[,c(1,6,7,11)])
```

```
## id_assessment pass code_module weight
## 1          1752 TRUE          AAA    10
## 2          1752 TRUE          AAA    10
## 3          1752 TRUE          AAA    10
## 4          1752 TRUE          AAA    10
## 5          1752 TRUE          AAA    10
## 6          1752 TRUE          AAA    10
```

#Final assessment average per student per module

```
avg_grade_others_exams <-
  student_group_others_exams %>%
  dplyr::group_by(id_student, code_module, code_presentation) %>%
  mutate(avg_grade = sum(weight_grade)) %>%
  select("id_student", "code_module", "code_presentation", "avg_grade")

head(avg_grade_others_exams)
```

```
## # A tibble: 6 × 4
## # Groups:   id_student, code_module, code_presentation [6]
## id_student code_module code_presentation avg_grade
##      <int> <chr>      <chr>          <dbl>
## 1    11391 AAA        2013J           82.4
## 2    28400 AAA        2013J           65.4
## 3    31604 AAA        2013J           76.3
## 4    32885 AAA        2013J           55
## 5    38053 AAA        2013J           66.9
## 6    45462 AAA        2013J           67.8
```

#Final exams scores

```
student_group_final_exams <-
  student_group_kpis %>%
  inner_join(final_exams, by = "id_assessment") %>%
  dplyr::rename("exams_score" = "score") %>%
  select("id_student", "code_module", "code_presentation", "exams_score")

head(student_group_final_exams)
```

```
## id_student code_module code_presentation exams_score
## 1    558914      CCC        2014B           32
## 2    559706      CCC        2014B           78
## 3    559770      CCC        2014B           54
## 4    560114      CCC        2014B           64
## 5    560311      CCC        2014B          100
## 6    560494      CCC        2014B           92
```

Having gathered the data from the assessments, let's check the data on student interactions with the university's virtual environment

Checking interactions:

The datasets pertaining to the university's virtual environment contain the student interaction feed with the available content. From this data, we can infer how a student was in touch with his subjects, whether he studied it solidly, and how he used the content.

```
#Reading the tables of interactions

student_vle <-
  read.csv(
    file = paste0(getwd(), "/Input/studentVle.csv")
  )

head(student_vle)
```

```
##   code_module code_presentation id_student id_site date sum_click
## 1      AAA      2013J      28400  546652  -10         4
## 2      AAA      2013J      28400  546652  -10         1
## 3      AAA      2013J      28400  546652  -10         1
## 4      AAA      2013J      28400  546614  -10        11
## 5      AAA      2013J      28400  546714  -10         1
## 6      AAA      2013J      28400  546652  -10         8
```

```
vle <-
  read.csv(
    file = paste0(getwd(), "/Input/vle.csv")
  )

head(vle)
```

```
##   id_site code_module code_presentation activity_type week_from week_to
## 1  546943      AAA      2013J      resource      NA      NA
## 2  546712      AAA      2013J      oucontent      NA      NA
## 3  546998      AAA      2013J      resource      NA      NA
## 4  546888      AAA      2013J      url          NA      NA
## 5  547035      AAA      2013J      resource      NA      NA
## 6  546614      AAA      2013J      homepage     NA      NA
```

If we look at the VLE table, we can identify that there are some data without reference to the period of use, so to make the analysis more feasible we will filter them out.

```
#Clearing the ELV data, because some samples do not have the reference week for the materials

vle <-
  vle %>%
  subset(!is.na(week_from))

head(vle)
```

```
##      id_site code_module code_presentation activity_type week_from week_to
## 114   546732      AAA      2013J      oucontent          2          2
## 199   546719      AAA      2013J      oucontent          1          1
## 211   546681      AAA      2013J      oucontent          1          1
## 265   877040      AAA      2014J      oucontent          2          2
## 324   877045      AAA      2014J      oucontent          1          1
## 392   877044      AAA      2014J      oucontent          1          1
```

Here we can track the average time after the start of the course that the student has taken to use the materials and the average number of clicks per material:

```
#Overall average per student per module
```

```
avg_per_student <-
  student_vle %>%
  dplyr::group_by(id_student, code_module, code_presentation) %>%
  mutate(
    date_mean = mean(date),
    sum_click_mean = mean(sum_click)) %>%
  select("id_student", "code_module", "code_presentation", "date_mean", "sum_click_mean")

head(avg_per_student)
```

```
## # A tibble: 6 × 5
## # Groups:   id_student, code_module, code_presentation [1]
##   id_student code_module code_presentation date_mean sum_click_mean
##      <int> <chr>      <chr>          <dbl>      <dbl>
## 1    28400 AAA      2013J          87.0        3.34
## 2    28400 AAA      2013J          87.0        3.34
## 3    28400 AAA      2013J          87.0        3.34
## 4    28400 AAA      2013J          87.0        3.34
## 5    28400 AAA      2013J          87.0        3.34
## 6    28400 AAA      2013J          87.0        3.34
```

Since we cannot identify performance faotres in the students who dropped out, we will take them out of the representative samples:

```
#Filtering only representative samples (According to the students' profile analysis)
```

```
representative_student_group_info <-
  student_info %>%
  subset(
    age_band == "0-35" &
    (
      highest_education == "A Level or Equivalent" |
      highest_education == "Lower Than A Level" |
      highest_education == "HE Qualification"
    ) &
    final_result != "Withdrawn"
  ) %>%
  distinct(id_student, .keep_all = TRUE)
```

```
#Compiling the relevant tables
```

```
df_1 <-  
  inner_join(avg_grade_others_exams, student_group_final_exams,  
             by = c("id_student", "code_module", "code_presentation"))  
  
df_2 <-  
  inner_join(representative_student_group_info, df_1,  
             by = c("id_student", "code_module", "code_presentation"))
```

```
## Warning in inner_join(representative_student_group_info, df_1, by = c("id_student", : Each  
row in `x` is expected to match at most 1 row in `y`.  
## i Row 3831 of `x` matches multiple rows.  
## i If multiple matches are expected, set `multiple = "all"` to silence this  
##   warning.
```

```
final_df <-  
  inner_join(df_2, avg_per_student,  
             by = c("id_student", "code_module", "code_presentation")) %>%  
  select(num_of_prev_attempts, final_result, avg_grade, exams_score, date_mean, sum_click_mean)
```

```
## Warning in inner_join(df_2, avg_per_student, by = c("id_student", "code_module", : Each row  
in `x` is expected to match at most 1 row in `y`.  
## i Row 1 of `x` matches multiple rows.  
## i If multiple matches are expected, set `multiple = "all"` to silence this  
##   warning.
```

```
head(final_df[, -2])
```

```
##   num_of_prev_attempts avg_grade exams_score date_mean sum_click_mean  
## 1                   0      89.65          94 119.3379      4.343939  
## 2                   0      89.65          94 119.3379      4.343939  
## 3                   0      89.65          94 119.3379      4.343939  
## 4                   0      89.65          94 119.3379      4.343939  
## 5                   0      89.65          94 119.3379      4.343939  
## 6                   0      89.65          94 119.3379      4.343939
```

```
summary(final_df[, -2])
```



```
## num_of_prev_attempts avg_grade exams_score date_mean
## Min. :0.0000 Min. : 3.72 Min. : 0.00 Min. : 27.01
## 1st Qu.:0.0000 1st Qu.:58.88 1st Qu.: 51.00 1st Qu.: 89.75
## Median :0.0000 Median :74.25 Median : 67.00 Median :103.28
## Mean :0.1059 Mean :70.98 Mean : 65.98 Mean :103.55
## 3rd Qu.:0.0000 3rd Qu.:86.15 3rd Qu.: 82.00 3rd Qu.:115.64
## Max. :5.0000 Max. :99.80 Max. :100.00 Max. :230.00
## NA's :20653
## sum_click_mean
## Min. : 1.077
## 1st Qu.: 2.273
## Median : 2.681
## Mean : 2.956
## 3rd Qu.: 3.331
## Max. :16.242
##
```

```
nrow(final_df[final_df$final_result == "Pass",])
```

```
## [1] 8043978
```

```
nrow(final_df[final_df$final_result == "Distinction",])
```

```
## [1] 2145211
```

```
nrow(final_df[final_df$final_result == "Fail",])
```

```
## [1] 1282021
```

With a much higher “Pass” count than the other labels, we should be on the lookout. Two outliers were detected: One with average clicks well above the standard values and another with a single occurrence from a number of previous attempts. To keep our data as consistent as possible, these cases will be removed.

```
final_df <-
  final_df %>%
  subset(sum_click_mean<10)
```

```
final_df <-
  final_df %>%
  subset(num_of_prev_attempts<4)
```

```
nrow(final_df)
```

```
## [1] 11435514
```

Separating the data to understand the profile of the students

who passed and those who failed

```
pass_student <-  
  final_df %>%  
    subset(final_result != "Fail")  
  
head(pass_student[, -2])
```

##	num_of_prev_attempts	avg_grade	exams_score	date_mean	sum_click_mean
## 1	0	89.65	94	119.3379	4.343939
## 2	0	89.65	94	119.3379	4.343939
## 3	0	89.65	94	119.3379	4.343939
## 4	0	89.65	94	119.3379	4.343939
## 5	0	89.65	94	119.3379	4.343939
## 6	0	89.65	94	119.3379	4.343939

```
nrow(pass_student)
```

```
## [1] 10155306
```

```
fail_student <-  
  final_df %>%  
    subset(final_result == "Fail")  
  
head(fail_student[, -2])
```

##	num_of_prev_attempts	avg_grade	exams_score	date_mean	sum_click_mean
## 42369	0	40.96	24	96.93333	5.114286
## 42370	0	40.96	24	96.93333	5.114286
## 42371	0	40.96	24	96.93333	5.114286
## 42372	0	40.96	24	96.93333	5.114286
## 42373	0	40.96	24	96.93333	5.114286
## 42374	0	40.96	24	96.93333	5.114286

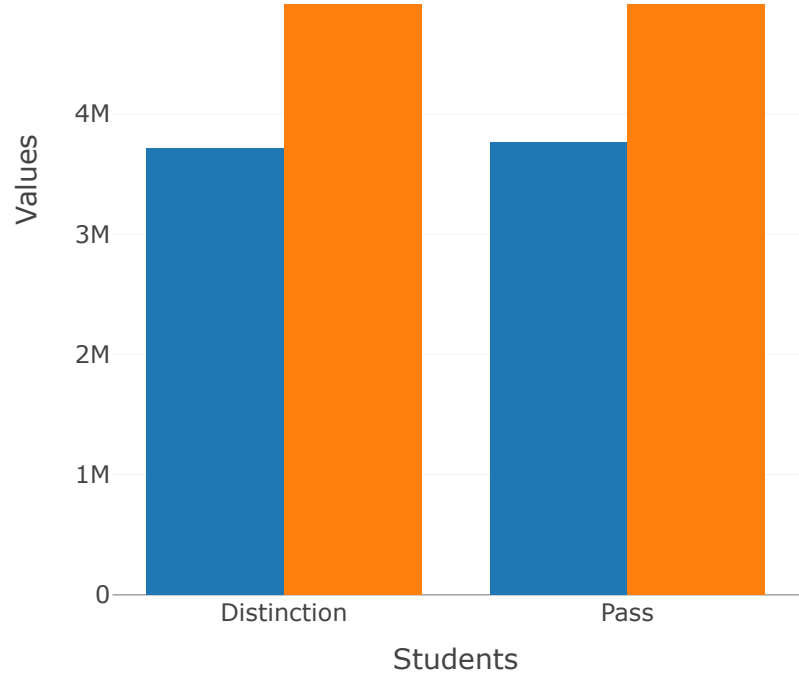
```
nrow(fail_student)
```

```
## [1] 1280208
```

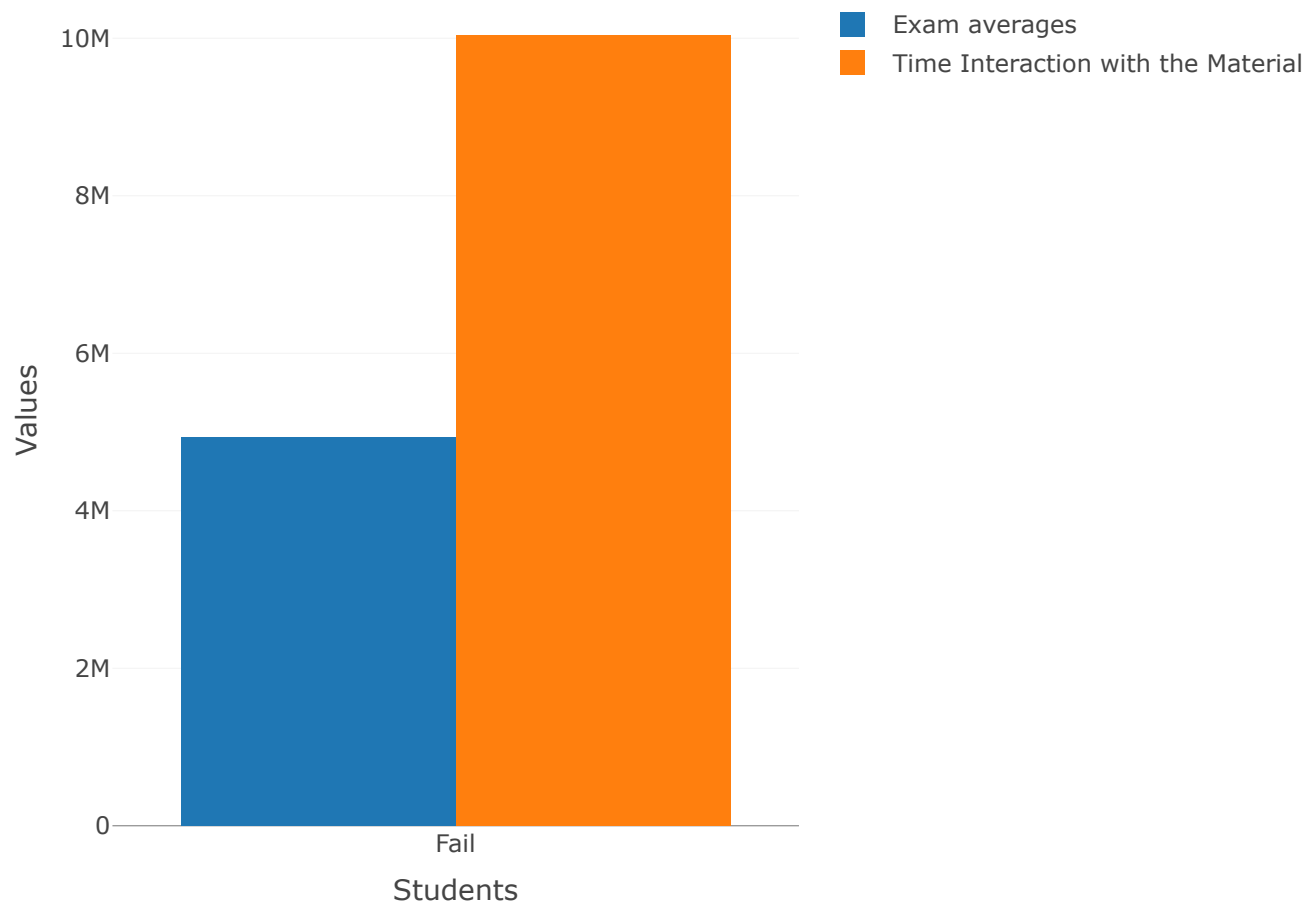
Since we have a large amount of samples, I will use below only the first one hundred thousand samples (100,000) to build the graphs:

Data from those who passed





Data of those who failed



6. Conclusion(III).

After the analysis performed, it can be seen that students who passed and performed better, overall had more interaction time with the material provided, as well as higher scores on the *Tutor Marked Assessment (TMA)* and *Computer Marked Assessment (CMA)* exams compared to those who did not pass. This may indicate that by focusing on initiatives to increase student interactivity with the online platform by increasing engagement

with the non-final exam papers, it is likely that the chance of success for students who did not pass would increase. For this analysis to become more assertive it would be necessary to follow up with some regression modeling and identify whether this hypothesis is valid.